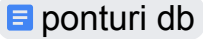


1. Ponturi db sem 3 -  ponturi db
2. Ce multi sunteti aici sa moara bibi
3. **ACID**
  - a. Atomicity: transactions are atomic (all or nothing)
  - b. Consistency: the consistency of the DB is preserved after executing a trans, all the constraints are still satisfied ( $\Leftrightarrow$  a trans is a correct program) the DB might be in a wrong state at some point during the trans, but it only has to be correct at the end
  - c. Isolation: a trans is protected from the effects of scheduling other trans concurrently
  - d. Durability: the changes of a successful trans will persist even if there's a crash before writing all changes to disk
    - i. Write-ahead log: changes are written to a log on the disk before being reflected in the DB
4. If 2 trans are just reading an object  $\Rightarrow$  no conflicts (exec order is not important)
5. If 2 trans are r/w different objects  $\Rightarrow$  no conflicts (...)
6. 2 trans r/w on the same object  $\Rightarrow$  possible conflicts (order is important)
  - a. WR (**dirty reads**, reading uncommitted data) - T2 reads sth written by T1
  - b. RW (**unrepeatable reads** - T2 writes sth already read by T1
  - c. WW (overwriting uncommitted data, *blind write*) - T2 writes sth already written by T1
7. **Serial schedules** - transactions are not interleaved
8. Serializable schedules iff their effect is identical to the effect of a serial schedule
9. Schedules are conflict equivalent ( $S1 \equiv_c S2$ , the order matters) iff  $\text{conflict}(S1) = \text{conflict}(S2) \Leftrightarrow S1, S2$  have the same op on the same objects + every pair of conflicting ops is ordered in the same way in  $S1$  and  $S2$
10.  $S$  is conflict serializable iff its precedence graph is acyclic
11. conflict serializable sch  $\equiv_c$  serializable sch
12. Conflict relation = set of pairs (op1, op2); op1 preceds op2 and in conflict
13. **Precedence graph**: one node for each committed trans in  $S$  + one arc  $T_i \rightarrow T_j$  if an op in  $T_i$  preceds and conflicts with an op in  $T_j$
14.  $S1, S2$  are **view equivalent** ( $S1 \equiv_v S2$ ) iff each transaction performs the same computation in  $S1$  and  $S2$  and  $S1, S2$  produce the same final DB state

15. Serial schedules  $\subset$  conflict serializable sch  $\subset$  view serializable sch  $\subset$  serializable sch  $\subset$  all schedules
16. Recoverable sch - T commits only after all the transactions whose changes T reads have committed
  - a. Avoid cascading aborts  $\Rightarrow$  recoverable sch
17. **Lock** - prevents a trans from accessing an object while another trans is accessing that object
18. Locks + transaction protocols  $\Rightarrow$  allow interleaved executions
19. Basically we can have multiple readers but just 1 writer at a time

Can you get a $\rightarrow$ lock if there is a $\downarrow$ lock	Shared lock	Exclusive lock
Shared lock	yes	no
Exclusive lock	no	no

20. Lock table - for each object, entries of the form (no. of trans with a lock on that obj, lock type, pointer to queue of lock requests)
21. Transactions table - for each trans, the list of locks held by it
22. Strict 2PL - acquires s/x lock before r/w; locks released when trans ends  $\Rightarrow$  only serializable sch
23. (Normal) 2PL - acquires...; once releasing a lock, it cannot request another one
24. Strict schedules - recoverable, no cascading aborts, if a trans is aborted, its ops can be undone;
  - a. Strict 2PL only allows strict sch
25. **Deadlock** - 2 trans wait for each other to free a resource
  - a. Usually the oldest trans has priority
  - b. Prevention policies  $\Rightarrow$  deadlocks cannot occur
    - i. **Wait-Die**: if T1 has higher priority, it waits, else is aborted
    - ii. **Wound-Wait**: if T1 has higher priority, T2 is aborted, else it waits
  - c. Waits-for graph: 1 node / active trans; arc T1 $\rightarrow$ T2 if T1 is waiting for T2 to release a lock; cycles  $\Rightarrow$  deadlocks; periodically checked by the DBMS
  - d. Timeout mechanism - if a trans waits too long for a lock to be released, a deadlock is assumed to exist  $\Rightarrow$  trans is terminated
26. Dirty writes are not allowed under any iso levels; X locks are always acquired for reading and always released at the end of the trans
27. **Read uncommitted**
  - a. A trans can read uncommitted data (dirty reads)
  - b. No S locks when reading

- c. Dirty reads, nonrepeatable reads, phantoms
- 28. **Read committed (default)**
  - a. Trans can only read committed data, but the read data can still be modified by another trans
  - b. S locks when reading, released asap
  - c. Nonrepeatable reads, phantoms
- 29. **Repeatable read**
  - a. Trans can only read committed data, the read data cannot be modified by another trans
  - b. S locks when reading, released at the end of the trans
  - c. Just phantoms
- 30. **Serializable**
  - a. RR + if the trans reads a set of objects based on a search predicate, that set cannot be modified by another trans (key-range locks)
  - b. Locks released at the end
  - c. No phantoms
- 31. **Snapshot - ?**
- 32. **Recovery manager** - ensures atomicity (uncommitted trans are undone) and durability (committed trans survive crashes)
- 33. Trans failure causes: system failure, app error, action by the transaction manager (ex. because of deadlock resolution), self-abort
- 34. **Steal** = the changes made by a trans can be written to disk before the commit (the BM wants to remove the frame of t1 to make space for a frame for t2 => t2 "steals" a frame from t1)
- 35. **No-steal** = changes cannot be written... => we don't have to undo changes in case of failure - but the BM needs to be able to hold all the changed pages, not always the case
- 36. **Force** = the changes are written to disk right when committing => no need to redo trans after crashes, because they have been written to disk - but if a page is committed many times by some transactions, it's written to disk that many times; no-force writes it only once
- 37. **No-force** = the changes might remain in the BP for a short while
- 38. Steal + no force -> most used
- 39. **ARIES** - steal + no force; after a crash, it does:
  - a. Analysis - determine active trans + dirty pages
  - b. Redo - reapply all changes, even those of aborted trans (bring the db to the state from when the crash occurred), in order
  - c. Undo - undo uncommitted trans, in reverse order

- d. Write-Ahead Logging - changes are first recorded in a log; the log is written to stable storage before the change is written to disk
- 40. **Log (journal)**
  - a. Log tail = most recent fragment; periodically forced to stable storage
  - b. LSN = log sequence number = id of a log record
  - c. Page LSN = the LSN of the most recent record describing a change to P
  - d. Log record = prev LSN, transID, type (update, commit etc)
  - e. A record is written for: update page, commit, abort, end, undo update (=> compensation log record CLR)
- 41. Transaction table - 1 entry / active trans
  - a. Trans ID, status (**in progress, committed, aborted**), lastLSN
- 42. Dirty page table - 1 entry / dirty page in BP
  - a. pageID, recLSN (LSN of log that first dirtied the page)
- 43. Checkpointing in ARIES
  - a. Write a begin\_checkpoint record
  - b. Write an end\_checkpoint record
  - c. After the end\_checkpoint is written to stable storage, write a master record to stable storage
- 44. Crash -> restart -> restore from the most recent checkpoint
- 45. **Analysis**
  - a. dirty page table gets a copy of the DirtyPageTable from the end checkpoint (most recent); same for trans table (TT)
  - b. end log record => remove the trans from TT
  - c. any other log record => add T to the TT if it's not there already, set its last LSN as the LSN of the crt log record;
  - d. if it's a commit => set T's status (in the TT) to C, otherwise U
  - e. if the page is not in DPT, add it; recLSN = crt LSN
  - f. otherwise just update recLSN
- 46. **Redo**
  - a. reapply the updates and CLRs, unless one of these conditions happen
    - i. the page (P) is not in DPT (usually their changes have already been written to disk)
    - ii. P in DPT, but P.recLSN > crtLSN
    - iii. P.pageLSN (from the actual DB on the disk) >= crtLSN
  - b. the actual redo: reapply the changes; P.pageLSN = crtLSN
- 47. **Undo**
  - a. loser transaction = active at the end of the crash

- b. toUndo = set of last LSNs of losers; empty => the undos have been completed
  - c. if L is a CLR, check undoNextLSN;
    - i. if not null => add undoNextLSN to toUndo
    - ii. else => do the redo?
  - d. if L is an update - write CLR, undo action, add L.prevLSN to toUndo
- 48. Security = protecting data against unauthorized users; integrity = .. authorized users
- 49. Access request = requested obj + requested operation + requesting user
- 50. Discretionary control - allowed operations are explicitly specified, anything else if forbidden
- 51. Mandatory control - each object has a classification level, each user has a clearance level
  - a. Bell and La Padula rules - x can retrieve y if  $\text{clearance}(x) \geq \text{classification}(y)$ ; x can update y if  $\text{clearance}(x) = \text{classification}(y)$
- 52. Unless specified otherwise, every command is a transaction
- 53. There exists an update lock (U) - deadlock avoidance mechanism
- 54. Key-range locking - lock existing data and data which doesn't exist (yet) based on a search predicate
- 55. SQL server uses deadlock detection; + the trans that is the least expensive to rollback is terminated
- 56. Algorithms for operators, based on
  - a. Iteration
  - b. Indexing
  - c. Partition
    - i. Sorting
    - ii. Hashing
- 57. **Access path** = way of retrieving tuples from a relation (file scan or index + condition)
  - a. Condition C matches index I if I can be used to retrieve just the tuples satisfying C
  - b. Hash index - used for equality selections; C has to have one term for each sk of I (C with <a, b> is not matched by I with <a, b, c>)
  - c. Tree-based index - used for all (<, >, ..), including equality; terms of C have to be a prefix of those of I
  - d. Most selective access path - retrieves the fewest pages (data retrieval cost is minimized)
  - e. General selection - the condition has to be in CNF; [L7, P16](#)

58. Joins

a. E (outer rel) - M pages;  $p_e$  records/ page; S = N pages;  $p_s$  records/ page

b. Iteration

i. **Simple nested loops join**; cost =  $M + p_e * M * N$  I/Os

```
foreach tuple e ∈ E do
    foreach tuple s ∈ S do
        if  $e_i == s_j$  then add <e, s> to the result
```

ii. **Page-Oriented nested loops join**; cost =  $M + M * N$

```
foreach page  $p_e$  ∈ E do
    foreach page  $p_s$  ∈ S do
        if  $e_i == s_j$  then add <e, s> to the result
```

iii. **Block nested loops join** - uses the buffer pages more effectively

```
foreach block  $b_e$  ∈ E do
    foreach page  $p_s$  ∈ S do
        {
            for all pairs of tuples <e, s> that meet the join
            condition, where  $e$  ∈  $b_e$  and  $s$  ∈  $p_s$ ,
            add <e, s> to the result
```

1. B pages in the buffer pool  $\Rightarrow$  B - 2 pages in a block (input / output each have 1 page)
2. Number of blocks =  $M / (B - 2)$  (rounded up)
3. Cost =  $M + \text{no. blocks} * N$

c. Indexing

i. Index nested loops join

```
foreach tuple e in E do
    foreach tuple s in S where  $e_i == s_j$ 
        add <e, s> to the result
```

1. Cost of examining the index on S = 2-4 for B+ tree index, 1.2 for hash index;
2. 1 extra I/O for reading a record in S, if clustered index; at most 1 if nonclustered
3. Cost:  $M + M * p_e * (\text{examine cost} + \text{read cost})$

d. Partitioning

i. Sort-merge join

1. Sort each relation (ex. with external merge sort)
2. Merge, with cost =  $M + N$  if the inner partition is scanned only once;  $M * N$  worst case
3. Cost = sorting costs + merge cost

ii. Hash join

1. Partitioning phase: 1 buffer page = input, B-1 pages = output; the hash function distributes tuples to the B-1 pages
2. **Partition** = collection of tuples with the same hash value, can be stored on 1+ pages when on disk
3. Probing phase: for each partition of the smaller relation, scan the other one for matching tuples; 1 page for input, 1 for output, B - 2 to read the partitions of the smaller relation
4. If a partition does not fit in memory => partition overflow => apply hash join recursively
5. Cost =  $3 * (M + N)$

59. If the data fits in main memory => internal sorting (ex. quicksort),  
else => external sorting

60. **Simple 2-way merge sort**

$$2 * N * ([\log_2 N] + 1) \text{ I/Os}$$

- a. Pass 0 -> read each page, sort it, save it back to disk => 1-page runs
- b. Pass 1+ -> use just 3 buffer pages, read + merge previous runs  
=> runs twice as long
- c. read + write each page at each pass (=> that first 2 I/O)

61. **External merge sort:** N = pages in input file; B = buffer pages =>  
cost =

$$2 * N * \left( \left[ \log_{B-1} \left\lceil \frac{N}{B} \right\rceil \right] + 1 \right) \text{ I/Os}$$

- a. Pass 0 -> read B pages at once, sort, write to disk

62. RLV (row-level versioning) - useful when we need committed data,  
but not necessarily the most recent version

- a. **XSN** = transaction sequence number; each row version is marked  
with ^ of the trans that changed the row

- b. The versions are kept in a linked list
  - c. **Read committed snapshot isolation** - operations see the most recent committed data as of the **beginning of their execution**; consistent reads at the command level
  - d. **Full snapshot isolation** - ...**beginning of their transaction**; consistent.. transaction level
  - e. Increased concurrency, but update operations are slower and we also 'use' a *tempdb*; also, there is still the issue of simultaneous writers (update conflict)
63. Projection, based on
- a. **Sorting** - select only the required columns, sort this => a temporary table, from which we remove adjacent duplicate tuples
    - i. Cost:  $M + T$  ( $T$  is  $O(M)$ ) + sort cost +  $T$
    - ii. Last  $T$  only if we remove duplicates
    - iii. Sort cost: see external merge sort / 2way merge, but it's  $O(M \log M)$
  - b. **Hashing**
    - i. Partitioning phase - 1 input page,  $B - 1$  output => hash into  $B - 1$  partitions, no unwanted fields
    - ii. Duplicate elimination phase - one partition at a time; build a new hash table for each to detect duplicates
    - iii. Cost:  $M + 2 * T$  ( $T$  = no. of partitions)
64. Intersection = join with condition of equality on all fields
65. Cross-product = join with no condition
66. Union, based on
- a. Sorting - sort the relations + merge them, eliminating duplicates
  - b. Hashing - partition the relations with a hash function; build another hash table for each partition for one of the relations; scan the other relation, while comparing with the values in the second hash table
67. Aggregate operations
- a. Without group by - maintain some *running info* about tuples (ex. current max, sum etc); cost = scan cost
  - b. With group by - sort the relation on the grouping attr, then scan the relation to get the aggregates for each group
68. **Optimizer** - generates plans for a query, tries to select the best one (but that's not always possible); optimizes one block at a time
- a. Convention - the outer relation is the left left child of the join op (in an algebra tree)



69. Statistics maintained by the DBMS - updated periodically (R = relation, I = index)
  - a. NTuples(R), NPages(R)
  - b. NKeys(I) (distinct key values), INPages(I) (page count or, if b+ tree index -> leaf pages); IHeight(I) (just for tree indexes); ILow(I), IHigh(I) (min/ max key value)
70. Estimating result sizes
  - a. **Reduction factor (RF)** - estimation of how much each term in a WHERE eliminates candidate rows
  - b. The estimation = product of the relation cardinalities \* product of RFs
  - c. Condition: column = value
    - i. Index on column =>  $RF = 1 / NKeys(I)$
    - ii. No index =>  $RF = 1 / 10$
  - d. Condition: column1 = column2
    - i. One index on each =>  $RF = 1 / MAX(nkeys(i1), nkeys(i2))$
    - ii. Just one index (on either column) =>  $RF = 1 / nkeys(i)$
    - iii. No index =>  $RF = 1 / 10$
  - e. Condition: column > value
    - i. Index =>  $RF = (ihigh(i) - value) / (ihigh - ilow)$
    - ii. No index => a random val < 0.5
  - f. Condition: column IN (value list) =>  $RF = (RF \text{ for column = value}) * \text{list length}$
  - g. Condition: NOT =>  $RF = 1 - RF(\text{condition})$
71. Enumeration of alternative plans
  - a. Query with 1 relation in FROM
    - i. No joins / cross-products
  - b. ... several relations in FROM
72. DMV = dynamic management view
73. Query fine-tuning:
  - a. Identify waits - sys.dm\_os\_wait\_stats
  - b. Correlate waits with queues - sys.dm\_os\_performance\_counters
  - c. Drill down to DB / file level - sys.dm\_io\_virtual\_file\_stats
  - d. Drill down to process level
74. **Distributed DB**
  - a. Distributed data independence - when writing a query, the user doesn't know / care that the data is distributed
  - b. Distributed transaction atomicity - a trans that operates on different sites is still atomic; again, the user doesn't know / care that the data is distributed

- c. **Homogeneous** - same DBMS at every site; **heterogeneous** - different DBMSs
- d. Fragmentation - split the data into fragments, stored across several sites
  - i. Horizontal - basically selection => disjoint subsets of rows; reconstruct -> union
  - ii. Vertical - basically projection => disjoint subsets of columns; reconstruct -> natural join
  - iii. Hybrid - combination of both ^
- e. Replication - same fragment in multiple sites => increased data availability (in case one location fails, we still have others) + faster queries (no communication costs)
  - i. **Synchronous** = all data copies are up to date before the trans commits
    - 1. **Voting** - when writing, a trans modifies a majority of its copies; when reading, a trans has to read at least (n - the value of the majority)
      - a. Ex: n = 10; T1 modifies 7 copies => T2 has to read 4, to guarantee that at least 1 was modified by T1
    - 2. **Read-any write-all** - when writing, a trans modifies all the copies; when reading, it has to read just 1
      - a. Better because reading is more frequent than writing, and when voting, reading is quite expensive
      - b. However, this method will require x locks from all the other sites, which could be slow
  - ii. **Asynchronous** = synchronized periodically => time periods when some copies are outdated -> the user will have to take data distribution into account
    - 1. **Peer to peer**
      - a. A hierarchy of copies, only master copies can be changed => all their successors are changed
      - b. preferred when there are no conflicts / data copies are disjoint
    - 2. **Primary site** - just one master copy, with secondary copies at other sites; change propagation in 2 steps:
      - a. Capture - ..all the changes made by a trans to the master copy

- i. Log based => a 'change data table' - which will contain only the update log records of committed trans
    - ii. Procedural - a procedure is automatically invoked (ex. trigger); uses snapshots of the primary copy
  - b. Apply - propagate changes to secondary copies; either the primary site continually sends the data (CDT / snapshot) to the secondaries, or the secondaries periodically request it
  - c. Log + continuous apply => minimum delay
  - d. Procedure + request => most flexible
- 75. Distributed query processing
  - a. There is a shipping cost - we transfer pages between sites
  - b. Nonjoin queries
  - c. Join queries
    - i. Fetch as needed - for each tuple in R, retrieve the corresponding tuples in A
    - ii. Ship to one site - ship the entire R to a location, do the join there
    - iii. Semijoin
      - 1. project R onto the join columns and send the projection to a location
      - 2. we do the join there ("reduction of A with respect to R")
      - 3. ship the reduction of A back to the first location, where we join R with the reduction of A
      - 4. Useful if there is a selection on one of the relations
    - iv. Bloomjoin
      - 1. Hash the tuples in R; use a bit vector to store the hash results, send the bit vector to another location
      - 2. Same hash but for A, discard the tuples whose hash value don't exist in the other bit vector => "*reduction of A with respect to R*", send the reduction to the first location
      - 3. Join R with the reduction of A
- 76. Distributed catalog management
  - a. Centralized system catalog - one site contains all the info; vulnerable to **single-site failures(SSF)**

- b. Global system catalog at each site - every site has all the info; no longer vulnerable to ssf<sup>^</sup>, but local autonomy is compromised (local changes have to be propagated to all other sites)
  - c. Local catalog at each site - each site has only local data; a catalog keeps track of all fragments / replicas of its relations - if we move / create a new replica => have to update the original catalog
- 77. Distributed transaction management
  - a. Distributed concurrency control
    - i. **Lock management**
      - 1. Centralized - one site does all the locking, vulnerable to SSF
      - 2. Primary copy
        - a. each obj has a primary copy stored at some site with some lock manager
        - b. lock requests of all copies of the obj are handled by that lock manager
        - c. not vulnerable to SSF
        - d. when reading another copy of the same obj, stored at a different site, we need to communicate with both sites
      - 3. Fully distributed - each obj has a copy...; but here the lock req are handled by the lock manager of the site where the copy is stored => no longer need to communicate with 2 sites
    - ii. **Distributed deadlocks**
      - 1. Each site has a local waits-for graph, but there could exist global deadlocks even if there are no cycles in the local graphs
      - 2. Centralized detection algorithm - all the local graphs are periodically sent to a single site, where they are united => global graph
      - 3. Hierarchical .. - there is a hierarchy of sites; each site sends the graph to its parent
      - 4. Based on timeout - trans is aborted if it waits too long; could be the only option in heterogeneous systems (different dbms => the graphs might not be compatible)

5. Phantom deadlock - a trans aborts before the global graph is updated => global thinks there is a deadlock and might abort another transaction

78. Distributed recovery

- a. (for a trans T) coordinator = transaction manager at the site where T originated; subordinates = .. sites where T's subtransactions execute
- b. **Two-phase commit protocol (2PC)**
  - i. 2 rounds of messages, both init by the coord (voting / termination)
  - ii. The user decides to commit T => commit command sent to T's coordinator => start 2PC
    1. Coord sends a "prepare" msg to each subordinate
    2. The subordinate decides whether to commit / abort and responds "yes" / "no"
    3. Coord receives "yes" => sends "commit" msg to all subord; ...at least a "no" or no msg at all from a subord => coord sends "abort"...
    4. Subord receives "abort" => writes an abort log record, responds with an "ack" msg, aborts the trans; ... receives "commit" => commit log, "ack" msg, commit trans
    5. Coord receives all "ack" msg => write end log
  - iii. Msg are sent only after the log has been forced to stable storage
- c. Restart after failure
  - i. If there is a commit / abort log => must redo / undo T; coord periodically sends commit / abort msg to subord until it receives all "ack" responses, then write and end log
  - ii. If there is a prepare log but no commit / abort => subord contact their coord until they get the trans status
  - iii. If there are no commit / abort / prepare logs => abort T, undo T
- d. Link and remote site failures: remote site R doesn't respond
  - i. If the crt site S is coord => S should abort T
  - ii. If S is a subord and hasn't voted yet(?) => S should abort T
  - iii. .. subord and voted "yes" => S blocked until T's coord responds